

Practical work 02 – 27/9/2018

Gradient Descent

Objectives

The main objectives of this Practical Work for Week 2 are the following :

- a) Get more experienced in python, particularly with numpy.
- b) Refresh or deepen your maths skills (multi-variate calculus)
- c) Implement gradient descent for the perceptron model and then apply it to MNIST.
Study some of the main features.

Submission

- **Deadline** : Wednesday 3 October, 12am
- **Format** :
 - Exercise 1 (Numpy)
 - iPython notebook `numpy-tutorial-stud.ipynb` completed with your solutions.
 - Exercise 2 (Sigmoid Function) :
 - pdf of your handwritten solutions.
 - iPython notebook with the sigmoid function and the plot (incl. derivative).
 - Exercise 3 (Gradient Descent Implementation) :
 - iPython notebook `MNIST_Light_Binary_PW-stud.ipynb` completed with your solutions.
 - Small pdf-report with plots and the answers to the questions.

Exercise 1 Numpy in a Nutshell

This exercise is to get you more familiar with `numpy`. Read the content of the ipython notebook `numpy-tutorial-stud.ipynb` that you will find on Moodle. Pay a special attention to the *broadcasting* section that allows to gain significant speedup when processing large numpy arrays. Regarding this, it is usually more efficient to use *broadcasting* instead of for loops in your code. At the end of the tutorial, you have to complete some manipulations of images stored by arrays.

Exercise 2 Sigmoid Function

- (a) Compute the derivative of the sigmoid function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- (b) Show that the derivative fullfills the equation

$$\sigma'(z) = \sigma(z) \cdot (1 - \sigma(z))$$

- (c) Compute the first and second derivative of

$$\zeta(z) = -\log(\sigma(-z)) \quad \log : \text{Natural Logarithm}$$

Compute the limits for $z \rightarrow \pm\infty$. Create a plot of ζ .

Remark : This function is also referred to as softplus, a smooth alternative of $x^+ = \max(0, x)$, the *rectifier*.

- (d) Implement the sigmoid function in a iPython Notebook. Make it work such that you can pass in numpy arrays of arbitrary shape and the function is applied element-wise. Plot the sigmoid function and its derivative by using matplotlib.
- (f) Show that the function

$$c_1(x) = (\sigma(x) - 1)^2$$

is non-convex. (Hint : Consider the second derivative.)

Explain in which situations (initial settings) optimising $c_1(x)$ with gradient descent may become difficult. For the explanation create a plot.

- (g) Compute the first and second derivative of the function

$$c_2(x) = -(y \log(\sigma(w \cdot x)) + (1 - y) \log(1 - \sigma(w \cdot x)))$$

with respect to $w \in \mathbb{R}$ and with $y \in \{0, 1\}$. Show that c_2 is convex.

Exercise 3 Gradient Descent for Perceptron

Implement gradient descent learning for the single perceptron and analyse the results. Do this on the basis of the iPython notebook `MNIST_Light_Binary_PW-stud.ipynb`. Do this by only using numpy functionality (scikit learn only for loading the data and splitting into train and test sets). The sections of code that you need to implement are marked again with

```
### START YOUR CODE ###
```

```
### END YOUR CODE ###
```

Proceed as follows :

- (a) Implement the sigmoid function, CE and MSE cost, associated update rules (`step_CE` and `step_MSE`) and the `optimise` function. Make sure that your implementation works also for larger images such as Original MNIST or other arbitrary datasets. This requires that you carefully design the shapes of the arrays used in the functions and the weights. (The shapes are indicated in the function descriptions.)
- (b) Plot the learning curves when optimising with the cross-entropy cost : Cost, Error Rate, Learning Speed and convince yourself that you obtain curves similar to as seen in the class.
- (c) Analyse the dependency of the final error rate on the number of epochs. What is the goal of the learning and how many epochs make sense? (Choose here the learning rate $\alpha = 0.5$.)
- (d) Analyse the dependency on the learning rate by trying different values, e.g.

0.01, 0.05, 0.1, 0.5, 1.0, 1.5, 2.0, 5.0, 10.0.

Determine the reasonable number of epochs to learn for each learning rate. Describe what you observe. How large can you choose the learning rate until the learning breaks down? Also check the learning speed (length of the gradient) and interpret what you see.

- (d) Analyse the dependency on the cost function. Compare the according error rates in a single plot. Give an explanation for why one of the cost functions works better.
- (e) Plot the weights finally obtained from learning as image and compare it with the misclassified test images. Try to explain.
- (g) Learn binary classification for all the digits (using reasonable values for the learning rate and the number of epochs). Compare the finally obtained error rates. Which digits work specifically well - which ones rather bad? Interpret these findings. Plot also the weights vector. Do you observe that for one of the problems the training set is linearly separable? Why?

Exercise 4 Optional : Digits Classification

The functions implemented in the previous exercise can now be expanded to build a digits classification system. Proceed as follows :

- a) For each digit separately, compute the probability that a given image depicts the digit.
- b) Then, pick the digit with largest probability and use this as the predicted digit.
- c) Finally, compute the error rate on the test dataset.

Use the functionality for binary classification implemented in Exercise 3.

Exercise 5 Optional : Review Questions

- a) Explain why normalisation is beneficial.
- b) In what sense are optimisation techniques important for machine learning problems?
- c) Describe in what problems gradient descent is applicable. In what problems will it necessarily lead to a unique solution? Describe what can go wrong in more general problems and why.
- d) Why is learning of classification problems with MSE cost considered more difficult?
- e) What may happen if the learning rate is chosen too large?
- f) Why is it possible that the test error starts increasing after some epochs of training?
- g) Is it becoming more or less difficult to reach small values of the cost function if we have more training data?
- h) Summarise the pro's and con's of BGD, SGD and MBGD.
- i) Assume you have a ML problem with
 - a model function that is very expensive to compute
 - lots of training data
 - a 32-cores machine but limited RAM.

What gradient descent method would you choose and why?

Exercise 6 Optional : Reading Assignment

Read the start of the Section 4.3 on *Gradient-Based Optimization* in the *Deep Learning* book by Ian Goodfellow et al (see <https://www.deeplearningbook.org/contents/numerical.html>). Students that have learned multi-variate calculus will find interesting reading also in Section 4.3.1.